



1. Dijkstra's Algorithm

Definition & Explanation

Dijkstra's algorithm finds the shortest path from a source node to all other nodes in a weighted graph with non-negative edge weights.

It uses a greedy approach with a priority queue (min-heap) to always expand the closest unvisited node.

Code Example

```
import heapq

def dijkstra(graph, start):
    # graph: adjacency list, e.g., {0: [(1, 2), (2, 4)], ...}
    heap = [(0, start)] # (distance, node)
    distances = {node: float('inf') for node in graph}
    distances[start] = 0
    visited = set()
    while heap:
        curr_dist, node = heapq.heappop(heap)
        if node in visited:
            continue
        visited.add(node)
        # Main Dijkstra step: update neighbors with shorter paths
        for neighbor, weight in graph[node]:
            if curr_dist + weight < distances[neighbor]:
                distances[neighbor] = curr_dist + weight
                heapq.heappush(heap, (distances[neighbor], neighbor))
    return distances
```

Example usage

```
graph = {  
    0: [(1, 4), (2, 1)],  
    1: [(3, 1)],  
    2: [(1, 2), (3, 5)],  
    3: []  
}  
  
print(dijkstra(graph, 0)) # Output: {0: 0, 1: 3, 2: 1, 3: 4}
```

Where is the algorithm applied?

- The main Dijkstra step is updating the shortest known distance to each neighbor and pushing it into the heap.
-

Big O Notation:

- **Time Complexity:** $O((V + E) \log V)$
Why?
 - Each node is pushed/popped from the heap at most once ($O(V)$ heap operations).
 - Each edge is checked once ($O(E)$).
 - Each heap operation is $O(\log V)$.
 - Total: $O((V + E) \log V)$.
 - **Space Complexity:** $O(V)$
Why?
 - The distances dictionary and heap store up to V entries.
-



2. Bellman-Ford Algorithm

Definition & Explanation

Bellman-Ford finds the shortest path from a source node to all other nodes in a weighted graph, even with negative edge weights (but no negative cycles).

It relaxes all edges $V-1$ times, where V is the number of vertices.

Code Example

```
def bellman_ford(graph, V, start):
    # graph: List of edges (u, v, w)
    distances = [float('inf')] * V
    distances[start] = 0
    # Relax all edges V-1 times
    for _ in range(V - 1):
        for u, v, w in graph:
            if distances[u] + w < distances[v]:
                distances[v] = distances[u] + w
    # Check for negative cycles
    for u, v, w in graph:
        if distances[u] + w < distances[v]:
            raise ValueError("Graph contains a negative-weight cycle")
    return distances

# Example usage
edges = [
    (0, 1, 4), (0, 2, 1),
    (2, 1, 2), (1, 3, 1),
    (2, 3, 5)
]

print(bellman_ford(edges, 4, 0)) # Output: [0, 3, 1, 4]
```

Where is the algorithm applied?

- The main Bellman-Ford step is relaxing all edges $V-1$ times.
-

Big O Notation:

- **Time Complexity:** $O(V \times E)$
Why?
 - For each of $V-1$ iterations, all E edges are checked and possibly relaxed.
 - **Space Complexity:** $O(V)$
Why?
 - The distances list stores one value per vertex.
-



3. Floyd-Warshall Algorithm

Definition & Explanation

Floyd-Warshall finds shortest paths between all pairs of nodes in a weighted graph (can handle negative weights, but not negative cycles).

It uses dynamic programming to update the shortest path between every pair using each node as an intermediate.

Code Example

```
def floyd_warshall(matrix):
    # matrix: adjacency matrix, matrix[i][j] = weight or float('inf')
    n = len(matrix)
    dist = [row[:] for row in matrix] # Copy of the matrix
    # Main Floyd-Warshall step: try every node as an intermediate
    for k in range(n):
        for i in range(n):
            for j in range(n):
                if dist[i][k] + dist[k][j] < dist[i][j]:
                    dist[i][j] = dist[i][k] + dist[k][j]
    return dist

# Example usage
INF = float('inf')
matrix = [
    [0, 3, INF, 7],
    [8, 0, 2, INF],
    [5, INF, 0, 1],
    [2, INF, INF, 0]
]
print(floyd_warshall(matrix)) # Output: Shortest path matrix for all pairs
```

Where is the algorithm applied?

- The main Floyd-Warshall step is updating $\text{dist}[i][j]$ using node k as an intermediate.
-

Big O Notation:

- **Time Complexity:** $O(V^3)$
Why?
 - Three nested loops over V nodes: for each pair (i, j) , try every possible intermediate node k .
 - **Space Complexity:** $O(V^2)$
Why?
 - The distance matrix stores shortest paths for every pair of nodes.
-

4. Union-Find (Disjoint Set Union, DSU)

Definition & Explanation

Union-Find is a data structure to efficiently manage a collection of disjoint sets, supporting:

- **Find:** Determine which set an element belongs to.
- **Union:** Merge two sets.

Used in Kruskal's MST, cycle detection, etc.

Code Example

```
class UnionFind:
    def __init__(self, n):
        self.parent = list(range(n))
        self.rank = [0] * n

    def find(self, x):
        # Path compression: flatten the tree
        if self.parent[x] != x:
            self.parent[x] = self.find(self.parent[x])
        return self.parent[x]

    def union(self, x, y):
        # Union by rank: attach smaller tree to larger
        rootX = self.find(x)
        rootY = self.find(y)
        if rootX == rootY:
            return False # Already connected
        if self.rank[rootX] < self.rank[rootY]:
            self.parent[rootX] = rootY
        elif self.rank[rootX] > self.rank[rootY]:
            self.parent[rootY] = rootX
        else:
            self.parent[rootY] = rootX
            self.rank[rootX] += 1
        return True

# Example usage
uf = UnionFind(5)
uf.union(0, 1)
uf.union(1, 2)
print(uf.find(2)) # Output: 0 (root of the set containing 2)
```

```
print(uf.find(3)) # Output: 3 (root of the set containing 3)
```

Where is the algorithm applied?

- The **find** method uses path compression, and **union** uses union by rank for efficiency.

Big O Notation:

- **Time Complexity:** $O(\alpha(n))$ per operation (almost constant)
Why?
 - With path compression and union by rank, the amortized time per operation is $O(\alpha(n))$, where α is the inverse Ackermann function (grows extremely slowly).
- **Space Complexity:** $O(n)$
Why?
 - The parent and rank arrays store one value per element.



5. Practice Problem

Kruskal's Minimum Spanning Tree (MST) using Union-Find

```
def kruskal(n, edges):
    # edges: List of (weight, u, v)
    uf = UnionFind(n)
    mst = []
    edges.sort() # Sort edges by weight
    for weight, u, v in edges:
        if uf.union(u, v):
            mst.append((u, v, weight))
    return mst
```



```
# Example usage
edges = [
    (1, 0, 1), (4, 0, 2), (3, 1, 2),
    (2, 1, 3), (5, 2, 3)
]
print(kruskal(4, edges)) # Output: MST edges
```

Where is the algorithm applied?





- Union-Find is used to check if adding an edge creates a cycle.
-

Big O Notation:

- **Time Complexity:** $O(E \log E + E \alpha(V))$
Why?
 - Sorting edges is $O(E \log E)$.
 - Each union/find is $O(\alpha(V))$, and there are up to E edges.
 - **Space Complexity:** $O(V)$
Why?
 - Union-Find data structure uses $O(V)$ space.
-



Summary Table

 Algorithm	 Time Complexity	 Space Complexity	 Why? (Time)
Dijkstra	$O((V + E) \log V)$	$O(V)$	Heap for closest node, check all edges
Bellman-Ford	$O(V \times E)$	$O(V)$	Relax all edges $V-1$ times
Floyd-Warshall	$O(V^3)$	$O(V^2)$	Triple nested loop for all pairs
Union-Find	$O(\alpha(n))$ per operation	$O(n)$	Path compression + union by rank
Kruskal's MST	$O(E \log E + E \alpha(V))$	$O(V)$	Sort edges, union-find for cycle detection