

Amazon Interviewer:

Hi Can, welcome back. We'll do a DSA problem framed in an Amazon fulfillment context. Please think aloud. Ready?

Candidate (Excellent):

Ready. I'll summarize the problem after you present it, then ask clarifying questions and proceed.

Interviewer — Problem Statement (Amazon context)

At an Amazon fulfillment center, a picker's handheld scanner records a sequence of scanned **SKU IDs** as they walk past bins. For a specific customer order, we also have a **requirement**: each SKU in the order and how many units are needed.

- You're given:
 - **scans**: an array of SKU IDs (strings or integers) in the order they were scanned.
 - **required**: a dictionary `sku -> quantity` (all quantities ≥ 1).

Task: Find the **shortest contiguous segment** of **scans** that contains **at least** the required quantity for **every** SKU in **required**.

Return the segment as **zero-based indices** `[start, end]`. If no such segment exists, return `[-1, -1]`.

Examples

1.

```
scans = ["A", "D", "B", "C", "A", "B", "C", "A"]
required = {"A":2, "B":1, "C":1}
```

Output: `[4, 7]`

Why: The segment `["A", "B", "C", "A"]` (indices 4..7) contains $A \times 2$, $B \times 1$, $C \times 1$, and its length is 4. Earlier valid segment `0..4` (`["A", "D", "B", "C", "A"]`) has length 5, so `[4, 7]` is shorter.

2.

```
scans = ["x", "y", "z"]
required = {"x":1, "y":1, "z":1}
```

Output: `[0, 2]` (the entire array)

3.

```
scans = ["p", "q"]
required = {"p":1, "q":2}
Output: [-1, -1] (not enough q)
```

Constraints

- $1 \leq \text{len}(\text{scans}) \leq 2 * 10^5$
- $1 \leq \text{len}(\text{required}) \leq 10^5$
- Each $\text{required}[\text{sku}] \geq 1$
- IDs are hashable (string or int)

Candidate — Problem Summary:

We're given a scan sequence and a multiset of required SKUs with counts. We need the **shortest contiguous subarray** that covers all required counts (\geq per SKU). If none exists, return $[-1, -1]$. Indices are zero-based. Constraints suggest we need around $O(n)$ or $O(n \log n)$; given the contiguity requirement, a **linear sliding window** sounds promising.

Did I capture this correctly?

Interviewer:

Perfect.

Candidate — Clarifying Questions:

1. If `required` is empty (no items needed), should I return a zero-length segment like $[0, -1]$ or $[-1, -1]$?
2. If multiple shortest segments exist, any one is acceptable?
3. `scans` can contain SKUs not in `required`; these are allowed inside the segment, correct?

Interviewer:

1. Treat `required` as non-empty for this interview (each quantity ≥ 1).

2. Yes, any shortest segment is fine.
 3. Correct—extra SKUs may appear; they don't invalidate the window.
-

Candidate — Approach Exploration

Brute force (reject):

Enumerate all subarrays and check coverage $\rightarrow O(n^2)$ windows and each check involves counting \rightarrow too slow for $2e5$.

Chosen Pattern (used because it fits and optimizes):

Pattern #1 — Sliding Window with two pointers and frequency maps.

Why it fits:

We're looking for a **shortest contiguous** segment that satisfies a **per-window constraint** ("window coverage \geq required counts"). Sliding window lets us expand to reach feasibility, then **shrink greedily** to minimal length.

How it works (with the example integrated):

Example 1

```
scans = ["A", "D", "B", "C", "A", "B", "C", "A"]
required = {"A":2, "B":1, "C":1}
```

Initialize state

- `left_index = 0`
- `best_start_index = -1, best_end_index = -1, best_window_length = +∞`
- `required_counts = {"A":2, "B":1, "C":1}`
- `required_types = 3` (distinct required SKUs)
- `window_counts = {}` (all 0)
- `satisfied_types = 0`

Expand `right_index` from 0 to end:

1. `right=0 \rightarrow "A"`
`window_counts["A"]=1 (<2) \rightarrow satisfied_types=0 \rightarrow not feasible.`

2. `right=1` → "D"
Not required → ignore for satisfaction → **not feasible**.
3. `right=2` → "B"
`window_counts["B"]`=1 meets B×1 → `satisfied_types`=1 → **not feasible** yet.
4. `right=3` → "C"
`window_counts["C"]`=1 meets C×1 → `satisfied_types`=2 → **not feasible** yet (A needs 2).
5. `right=4` → "A"
`window_counts["A"]`=2 meets A×2 → `satisfied_types`=3 (= `required_types`) → **feasible window [0..4]**.
Shrink from left to minimize:
 - `left=0` is "A" → dropping would make A 2→1 (<2) ⇒ **stop shrinking**.
Record best: `[0, 4]` (length 5).
6. `right=5` → "B"
 - `window_counts["B"]`=2
Still feasible. Try shrinking:
 - `left=0` is "A" → A would become 2→1 ⇒ **break feasibility** ⇒ stop.
Best remains `[0, 4]`.
7. `right=6` → "C"
 - `window_counts["C"]`=2
Still feasible. Try shrinking:
 - `left=0` "A" → would break ⇒ stop.
Best unchanged.
8. `right=7` → "A"
`window_counts["A"]`=3. **Shrink aggressively** while staying feasible:
 - Drop `left=0` "A": A 3→2 (still ≥2) ⇒ feasible → `left=1`
 - Drop `left=1` "D": not required ⇒ feasible → `left=2`
 - Drop `left=2` "B": B 2→1 (still ≥1) ⇒ feasible → `left=3`
 - Drop `left=3` "C": C 2→1 (still ≥1) ⇒ feasible → `left=4`
 - Next would be dropping `left=4` "A": A 2→1 (<2) ⇒ **stop**.
Now the minimal feasible window ending at `right=7` is `[4, 7]` (length 4).

Update best to [4, 7].

Final answer: [4, 7].

Key idea: expand **right** to become feasible, then move **left** to trim extraneous items while keeping the window feasible, updating the best answer whenever we're feasible.

Candidate — Pseudocode

```
required_counts = frequency map from required (sku -> needed_count)
need_kinds      = number of keys in required_counts
window_counts   = empty map (sku -> count in window)
covered_kinds   = 0
left            = 0
best_start, best_end, best_length = -1, -1, +inf

for right in range(len(scans)):
    sku = scans[right]
    add sku to window_counts
    if sku in required and window_counts[sku] just reached required_counts[sku]:
        covered_kinds += 1

    while covered_kinds == need_kinds:
        # Feasible: try to shrink from the left
        if right - left + 1 < best_length:
            best_length = right - left + 1
            best_start, best_end = left, right

        left_sku = scans[left]
        remove left_sku from window_counts
        if left_sku in required and window_counts[left_sku] just dropped below
required_counts[left_sku]:
            covered_kinds -= 1
        left += 1

return [best_start, best_end]
```

Pattern Callout: Using **Sliding Window (#1)** because we need a shortest **contiguous** segment satisfying per-window frequency constraints.

Interviewer:

Looks good. Please implement with meaningful variable names and explanatory comments.

Candidate — Code (Python, descriptive names + comments)

```
from collections import Counter
from math import inf
from typing import List, Dict, Tuple, Hashable

def shortest_cover_segment(scans: List[Hashable], required: Dict[Hashable, int])
-> List[int]:
    """
        Find the shortest [start, end] segment of `scans` that covers all required
        SKUs with
        at least the specified counts. Return [-1, -1] if not possible.

        Sliding Window (expand right, shrink left) with frequency maps.

        Time: O(n) average, each index enters/leaves the window at most once.
        Space: O(U) where U is number of distinct SKUs seen in the window.
    """
    if not scans or not required:
        # Per interview, required is non-empty; returning [-1,-1] for safety if
        empty.
        return [-1, -1]

    required_counts: Dict[Hashable, int] = dict(required)
    needed_kinds: int = len(required_counts)

    window_counts: Counter = Counter()
    covered_kinds: int = 0

    best_start: int = -1
    best_end: int = -1
    best_length: int = float('inf')

    left: int = 0

    for right, sku in enumerate(scans):
        # include current sku in the window
```

```

    window_counts[sku] += 1

    # If this sku is required and we just satisfied its required count,
    increment covered_kinds
    if sku in required_counts and window_counts[sku] == required_counts[sku]:
        covered_kinds += 1

    # If all required kinds are covered, try to shrink from the left to get a
    minimal window
    while covered_kinds == needed_kinds and left <= right:
        current_window_length = right - left + 1
        if current_window_length < best_length:
            best_length = current_window_length
            best_start = left
            best_end = right

        left_sku = scans[left]
        window_counts[left_sku] -= 1

        # If left_sku is required and we just fell below its needed count, we
        lose coverage
        if left_sku in required_counts and window_counts[left_sku] <
        required_counts[left_sku]:
            covered_kinds -= 1

        # Move left pointer forward to continue shrinking
        left += 1

    return [best_start, best_end]

```

Interviewer:

Great. Let's dry-run the primary example and then run a quick test harness.

Candidate — Dry Run (Example 1)

```

scans = ["A", "D", "B", "C", "A", "B", "C", "A"]
required = {"A":2, "B":1, "C":1}

```

- Expand until `right=4` (second "A"): now A×2, B×1, C×1 → feasible at `[0..4]` (length 5).

- Shrink fails immediately because dropping the left “A” breaks A×2. Best is [0, 4].
- Continue expanding to the end; each time we regain plenty of A’s, shrinking removes non-required “D” and older items.
- Final minimal window discovered: [4, 7] of length 4 ([“A”, “B”, “C”, “A”]).

Result: [4, 7].

Candidate — Tests (expected outputs in comments)

```
def _tests():
    print(shortest_cover_segment(
        ["A", "D", "B", "C", "A", "B", "C", "A"],
        {"A":2, "B":1, "C":1})) # [4, 7]

    print(shortest_cover_segment(
        ["x", "y", "z"],
        {"x":1, "y":1, "z":1})) # [0, 2]

    print(shortest_cover_segment(
        ["p", "q"],
        {"p":1, "q":2})) # [-1, -1]

    print(shortest_cover_segment(
        ["a", "b", "a", "b", "c"],
        {"a":1, "c":1})) # [2, 4] or [0,4], shortest is [2,4]

    print(shortest_cover_segment(
        ["a", "a", "b", "b", "c", "c"],
        {"a":2, "b":2, "c":2})) # [0, 5] (entire array)

# _tests()
```

Interviewer:

Explain your time and space complexity with reasoning, not just big-O symbols.

Candidate — Complexity (with reasoning)

- **Time Complexity:**

The right pointer advances from 0 to $n-1$ (n steps). The left pointer only moves forward and never resets—each index is removed from the window at most once. The inner `while` loop therefore advances `left` at most n total times across the run. All map updates and comparisons are $O(1)$ average (hash map).

Total: $\sim O(n)$ average time.

- **Space Complexity:**

`required_counts` holds at most the number of distinct required SKUs; `window_counts` holds at most the number of distinct SKUs present in the current window. In the worst case, this is $O(U)$, where $U \leq n$.

Total extra space: $O(U)$ (commonly summarized as $O(n)$ in the worst case, but practically bounded by the unique SKUs).

Interviewer (targeted follow-ups):

1. What if some `required` counts are larger than the total occurrences in `scans`?
2. How do you ensure we always return the **shortest** feasible window?
3. If we also needed the **count** of shortest windows (how many distinct minimal segments), how might you extend this?

Candidate:

1. The algorithm will simply **never reach** `covered_kinds == needed_kinds`, so `best_start` stays `-1` and we return `[-1, -1]`.
2. We update the best answer **only** when the window is **currently feasible**, and then we **shrink greedily** from the left as far as possible before expanding again. That produces the minimal window for each `right`, and we track the global minimum over time.
3. To count minimal segments: when we shrink a feasible window, if the new window length equals the best length, increment a counter; if the new window is shorter, reset the counter to 1 and update the best. Be careful to distinguish equal-length windows starting at different positions.

Interviewer — Wrap-up:

Excellent session. You led with a clear summary, asked precise clarifications, chose **Sliding Window (Pattern #1)** because it truly fit, articulated the feasibility/shrinking invariant, wrote clean code with meaningful variable names and comments, and provided dry runs, tests, and well-reasoned complexity. This meets Amazon's bar.

Ratings

- Coding: **4/4**
- Problem Solving: **4/4**
- Communication: **4/4**