

Amazon Interviewer:

Hi Can, welcome back. Let's do a DSA round in an Amazon operations context. I'll present the problem, then I'd like you to summarize it, ask clarifying questions, and think aloud as you solve it.

Problem — Minimum Average Processing Time at a Sort Center

At an Amazon sort center, each customer order turns into a **job** when it arrives. We have **one packing station** (single machine), jobs are **non-preemptive** (once started, you must finish), and you may pick **any available job** when the station is free.

For each job i :

- $arrival[i]$ — time it enters the system,
- $duration[i]$ — uninterrupted processing time the station needs.

Goal: choose a processing order that **minimizes the average turnaround time** (also called average completion delay), defined for job i as:

$$turnaround(i) = completion_time(i) - arrival[i]$$

We return the **floor** of the average over all jobs.

If the station is idle and **no job has arrived yet**, it **waits** until the next arrival.

Return: an integer = $\text{floor}(\text{average}(turnaround(i)))$.

Constraints

- $1 \leq n \leq 2 * 10^5$
- $0 \leq \text{arrival}[i], \text{duration}[i] \leq 10^9$
- Single machine; non-preemptive

Worked Examples (with explanations)

Example 1

`arrival = [0, 1, 2]`

`duration = [3, 9, 6]`

Optimal schedule (Shortest-Job-First among available):

- $t=0$: only job0 (dur 3) is available \rightarrow run job0, completes at $t=3$, turnaround = $3 - 0 = 3$.
- $t=3$: jobs1(dur 9, arr 1), job2(dur 6, arr 2) are available \rightarrow pick shorter job2 \rightarrow completes at $t=9$, turnaround = $9 - 2 = 7$.
- $t=9$: run job1 \rightarrow completes at $t=18$, turnaround = $18 - 1 = 17$.

Sum = $3 + 7 + 17 = 27$, average = $27 / 3 = 9$, **answer = 9**.

Example 2

`arrival = [0, 0, 0]`

`duration = [2, 2, 2]`

All three jobs arrive at time 0 and take 2 time units each. Any order yields:

- First completes at $t=2 \rightarrow 2 - 0 = 2$
- Second completes at $t=4 \rightarrow 4 - 0 = 4$
- Third completes at $t=6 \rightarrow 6 - 0 = 6$
Sum = $2 + 4 + 6 = 12$, average = $12 / 3 = 4$, **answer = 4**.

Note: We're using **turnaround time** (completion – arrival). If you meant *queue waiting only*, we'd define it differently. For this round we stick to turnaround, as above.

Candidate (Excellent) — Problem Summary:

We have n non-preemptive jobs on a single machine. Each job has an arrival and a processing duration. When the machine is free, it can pick **any** job whose arrival time has passed. We want to minimize the **average turnaround** = completion – arrival, and return its floor. If no jobs are available, time advances to the next arrival. Constraints push us to an $O(n \log n)$ -type solution.

Did I capture that correctly?

Interviewer:

Perfect.

Candidate — Clarifying Questions:

1. If multiple jobs are available at the same time, I'm free to choose any ordering?
 2. Ties on duration don't matter for correctness, right?
 3. Confirm we return `floor(total_turnaround / n)` (integer division)?
-

Interviewer:

1. Yes, any available job.
 2. Correct; any tie-break among equal durations is fine.
 3. Yes, floor of the average.
-

Candidate — Approach Exploration

Brute-force (reject): Trying all valid schedules is factorial and intractable.

Greedy insight: To minimize average turnaround (equivalently, sum of completion minus arrival), when jobs are available, always process the **shortest duration next**. This is the classic **Shortest Job First (SJF)** strategy.

Pattern (used because it fits and optimizes):

- **Greedy Algorithms (#27)** — locally optimal choice (shortest available job) leads to global optimum.

- Backed by an **exchange argument**: if you ever run a longer job before a shorter one that's already available, swapping them does not increase (and typically reduces) the total completion delay.

Data structure: A min-heap (priority queue) keyed by **duration** to select the next shortest job among those that have arrived.

Now I'll turn this into a **concrete step-by-step plan** using **Example 1** (**arrival**=[0,1,2], **duration**=[3,9,6]):

1. **(We can skip a separate sort)** because **arrival** is already ascending.
→ *State now:* **arrival**[**i**], **duration**[**i**] represent job **i**, in arrival order.
2. **Maintain **time** and a pointer **i** over the arrays.**
→ Initialize: **time** = 0, **i** = 0, min-heap **pq** = [], **total_turnaround** = 0.
3. **Push all jobs with **arrival**[**i**] ≤ **time** into the min-heap (**duration**[**i**], **arrival**[**i**]).**
 - At start: **time**=0, **arrival**[0]=0 ≤ **time** → push (3, 0).
 - Heap now: [(3, 0)].
4. **If heap is empty, jump **time** to **arrival**[**i**].**
 - *Explanation:* This models idling. If no arrived jobs are available, we fast-forward to the next arrival to avoid spinning.
5. **Pop the shortest job; update time and turnaround.**
 - Pop (3, 0) → run → **time** = 0 + 3 = 3

- Add turnaround: $\text{time} - \text{arrival} = 3 - 0 = 3 \rightarrow \text{total} = 3$.

6. Repeat until all jobs processed.

- Now $\text{time}=3$, enqueue any arrivals ≤ 3 : jobs at $i=1$ ($\text{arr}=1, \text{dur}=9$) and $i=2$ ($\text{arr}=2, \text{dur}=6$) \rightarrow push $(9, 1)$ and $(6, 2) \rightarrow$ heap = $[(6, 2), (9, 1)]$
- Pop shortest $(6, 2) \rightarrow \text{time}=9$, add $9-2=7 \rightarrow \text{total} = 10$
- Pop $(9, 1) \rightarrow \text{time}=18$, add $18-1=17 \rightarrow \text{total} = 27$

7. Return $\text{total_turnaround} // n$.

- $27 // 3 = 9 \rightarrow$ matches Example 1.

This example-driven plan clarifies exactly what happens at each step.

This is $O(n \log n)$ due to sorting and heap operations.

Interviewer (Probe):

Give me the short correctness sketch for SJF here.

Candidate — Correctness (exchange argument, brief):

Consider two available jobs **A** and **B** with durations $a \leq b$. If we schedule **B** then **A**, their contributions to total turnaround are larger or equal compared to **A** then **B**. Swapping any such inversions repeatedly yields a schedule where shorter jobs are always processed before longer ones among available jobs—i.e., SJF—without increasing the total turnaround. Hence SJF is optimal.

Candidate — Pseudocode (concise)

```
jobs = sort by arrival asc
i = 0
time = 0
total = 0
heap = min-heap of (duration, arrival)

while i < n or heap not empty:
    if heap empty and time < jobs[i].arrival:
        time = jobs[i].arrival
    while i < n and jobs[i].arrival <= time:
        push (jobs[i].duration, jobs[i].arrival) to heap
        i += 1
    d, a = pop heap
    time += d
    total += time - a

return total // n
```

Interviewer:

Please implement it and then we'll dry-run both examples.

Candidate — Code (Python)

```
import heapq
from typing import List

def min_average_turnaround(arrival: List[int], duration: List[int]) -> int:
    """
    Compute floor(average turnaround) on a single non-preemptive machine.

    Turnaround(i) = completion_time(i) - arrival[i]
    Strategy: Shortest-Job-First among jobs that have arrived.
               Use a min-heap keyed by duration to pick the next job.

    Assumptions for this interview:
    - arrival is already sorted ascending.
    - arrival and duration are aligned by index (job i is (arrival[i], duration[i])).
    """

    n = len(arrival)
```



```

if n == 0:
    return 0 # no jobs => average turnaround is 0 by convention

i = 0 # pointer over (arrival, duration)
time = 0 # current machine time
total = 0 # sum of all turnaround times
pq = [] # min-heap of (duration, arrival) for available jobs

# Process until all jobs are scheduled (i == n) and heap is empty
while i < n or pq:
    # If no available jobs, jump time forward to the next arrival
    if not pq and i < n and time < arrival[i]:
        time = arrival[i] # idle → fast-forward to next job arrival

    # Enqueue every job that has arrived by 'time'
    # Since arrival is sorted, this loop advances 'i' monotonically.
    while i < n and arrival[i] <= time:
        heapq.heappush(pq, (duration[i], arrival[i]))
        i += 1

    # If we have any available job, run the shortest duration next.
    if pq:
        d, a = heapq.heappop(pq) # pick the job with smallest duration
        time += d # advance time by its duration
        total += (time - a) # add its turnaround = completion - arrival

```

```
# Return the floored average turnaround
return total // n
```

Interviewer:

Dry-run Example 1 and Example 2 exactly as stated earlier, and explain the outputs.

Candidate — Dry Runs**Example 1**

```
arrival = [0, 1, 2], duration = [3, 9, 6]
```

Sorted jobs: (0,3), (1,9), (2,6)

- $t=0$: push (3,0) \rightarrow pop (3,0), $t=3$, turnaround $3-0=3$, total=3
- $t=3$: push (9,1) and (6,2) \rightarrow pop (6,2), $t=9$, turnaround $9-2=7$, total=10
- $t=9$: pop (9,1), $t=18$, turnaround $18-1=17$, total=27
Average = $27/3 = 9 \rightarrow 9$

Example 2

```
arrival = [0, 0, 0], duration = [2, 2, 2]
```

- $t=0$: push all three $(2, 0)$
- Pop first: $t=2$, turnaround $2-0=2$, total=2
- Pop second: $t=4$, turnaround $4-0=4$, total=6
- Pop third: $t=6$, turnaround $6-0=6$, total=12
Average = $12/3 = 4 \rightarrow 4$

Interviewer:

Good. Add a quick test harness and list expected outputs.

Candidate — Tests (expected outputs in comments)

```
def _test():
    print(min_average_turnaround([0, 1, 2], [3, 9, 6]))    # 9
    print(min_average_turnaround([0, 0, 0], [2, 2, 2]))    # 4
    print(min_average_turnaround([5], [7]))                # 7 (finishes at 12)
    print(min_average_turnaround([0, 10, 20], [10, 1, 1])) # 4
    print(min_average_turnaround([0, 2, 4], [100, 2, 2])) # 64
# _test()
```

Interviewer:

Great. Can you give me the complexity with reasoning?

Candidate — Complexity (with reasoning)

- **Time:**

- We push each job **once** and pop **once** from the min-heap → $2n$ heap ops total, each $O(\log n) \rightarrow O(n \log n)$.
- The inner enqueue loop advances pointer i monotonically; the outer loop doesn't revisit work.
- **Overall:** $O(n \log n)$.

- **Space:**

- Heap may hold up to n jobs in the worst case; plus a few scalars.
 - **Overall extra space:** $O(n)$.
-

Interviewer (Targeted Follow-ups):

1. What happens with long idle gaps?
2. Are there any pitfalls with very large times or durations?

3. If we wanted the **actual schedule**, how would you return it?
-

Candidate:

1. When the heap is empty and there are still jobs left, I **jump time** to the next job's arrival. That prevents spinning and models idling accurately.
 2. Large values are fine in Python (unbounded ints). Complexity still holds since it depends on **n**, not on numeric magnitudes.
 3. I'd store **(duration, arrival, index)** in the heap and append **index** to a **schedule** list each time I pop. Return that list in addition to the average.
-

Interviewer — Wrap-up:

Excellent. You summarized first, asked precise clarifications, selected a **Greedy + Min-Heap** approach with a clear correctness sketch, wrote clean code, performed transparent dry runs (including detailed math for the 9 and 4 results), provided tests and a solid complexity explanation. This meets Amazon's bar.

Ratings

- Coding: **4/4**
- Problem Solving: **4/4**
- Communication: **4/4**

